

```
!pip install shap pyDOE2
from IPython.core.display import display, HTML
import regex as re
import lightgbm
import pandas as pd
import shap
import sklearn

import xgboost as xgb
from sklearn.model_selection import train_test_split
import lightgbm as lgb
```

```
↔ Requirement already satisfied: shap in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pyDOE2 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages
```

Patch to match style consistency

```
import matplotlib
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

from shap import Explanation
from shap.utils import format_value
from shap.plots import colors
from shap.plots._labels import labels
```

```
plt.rcParams['figure.dpi'] = 300
```

```
# TODO: If we make a JS version of this plot then we could let users click on a bar
# plot that is associated with that feature get overlaid on the plot...it would query
# why a feature is pushing down or up. Perhaps the best way to do this would be with
# of the bar...
```

```
def patch_waterfall(shap_values, max_display=10, show=True):
    """Plots an explanation of a single prediction as a waterfall plot.
```

The SHAP value of a feature represents the impact of the evidence provided by that feature on the model output. The waterfall plot is designed to visually display how the SHAP value for a feature moves the model output from our prior expectation under the background data distribution to the current prediction given the evidence of all the features.

Features are sorted by the magnitude of their SHAP values with the smallest magnitude features grouped together at the bottom of the plot when the number of features in the models exceeds the ``max_display`` parameter.

Parameters

shap_values : Explanation

A one-dimensional :class:`.Explanation` object that contains the feature importance values.

max_display : str

The maximum number of features to plot (default is 10).

show : bool

Whether ``matplotlib.pyplot.show()`` is called before returning.

Setting this to ``False`` allows the plot to be customized further after it has been created, returning the current axis via plt.gca().

Examples

See `waterfall plot examples` <https://shap.readthedocs.io/en/latest/example_notebooks/plotting/waterfall.html>

"""

```
# Turn off interactive plot
```

```
#if show is False:
```

```
#    plt.ioff()
```

```
# make sure the input is an Explanation object
```

```
if not isinstance(shap_values, Explanation):
```

```
    emsg = (
```

```

        "The waterfall plot requires an `Explanation` object as the "
        "`shap_values` argument."
    )
    raise TypeError(msg)

# make sure we only have a single explanation to plot
sv_shape = shap_values.shape
if len(sv_shape) != 1:
    msg = (
        "The waterfall plot can currently only plot a single explanation, but "
        f"matrix of explanations (shape {sv_shape}) was passed! Perhaps try "
        "`shap.plots.waterfall(shap_values[0])` or for multi-output models, "
        "try `shap.plots.waterfall(shap_values[0, 0])`."
    )
    raise ValueError(msg)

base_values = float(shap_values.base_values)
features = shap_values.display_data if shap_values.display_data is not None else
feature_names = shap_values.feature_names
lower_bounds = getattr(shap_values, "lower_bounds", None)
upper_bounds = getattr(shap_values, "upper_bounds", None)
values = shap_values.values

# unwrap pandas series
if isinstance(features, pd.Series):
    if feature_names is None:
        feature_names = list(features.index)
    features = features.values

# fallback feature names
if feature_names is None:
    feature_names = np.array([labels['FEATURE'] % str(i) for i in range(len(values))])

# init variables we use for tracking the plot locations
num_features = min(max_display, len(values))
row_height = 0.5
rng = range(num_features - 1, -1, -1)
order = np.argsort(-np.abs(values))
pos_lefts = []
pos_inds = []
pos_widths = []
pos_low = []
pos_high = []
neg_lefts = []
neg_inds = []

```

```

neg_widths = []
neg_low = []
neg_high = []
loc = base_values + values.sum()
yticklabels = [""] for _ in range(num_features + 1)

# size the plot based on how many features we are plotting
plt.gcf().set_size_inches(8, num_features * row_height + 1.5)

# see how many individual (vs. grouped at the end) features we are plotting
if num_features == len(values):
    num_individual = num_features
else:
    num_individual = num_features - 1

# compute the locations of the individual features and plot the dashed connect
for i in range(num_individual):
    sval = values[order[i]]
    loc -= sval
    if sval >= 0:
        pos_inds.append(rng[i])
        pos_widths.append(sval)
        if lower_bounds is not None:
            pos_low.append(lower_bounds[order[i]])
            pos_high.append(upper_bounds[order[i]])
        pos_lefts.append(loc)
    else:
        neg_inds.append(rng[i])
        neg_widths.append(sval)
        if lower_bounds is not None:
            neg_low.append(lower_bounds[order[i]])
            neg_high.append(upper_bounds[order[i]])
        neg_lefts.append(loc)
if num_individual != num_features or i + 4 < num_individual:
    plt.plot([loc, loc], [rng[i] - 1 - 0.4, rng[i] + 0.4],
             color="#bbbbbb", linestyle="--", linewidth=0.5, zorder=-1)
if features is None:
    yticklabels[rng[i]] = feature_names[order[i]]
else:
    if np.issubdtype(type(features[order[i]]), np.number):
        yticklabels[rng[i]] = format_value(float(features[order[i]]), "%0
    else:
        yticklabels[rng[i]] = str(features[order[i]]) + " = " + str(featu

# add a last grouped feature to represent the impact of all the features we d

```

```

if num_features < len(values):
    yticklabels[0] = "%d other features" % (len(values) - num_features + 1)
    remaining_impact = base_values - loc
    if remaining_impact < 0:
        pos_inds.append(0)
        pos_widths.append(-remaining_impact)
        pos_lefts.append(loc + remaining_impact)
    else:
        neg_inds.append(0)
        neg_widths.append(-remaining_impact)
        neg_lefts.append(loc + remaining_impact)

points = pos_lefts + list(np.array(pos_lefts) + np.array(pos_widths)) + neg_l
    list(np.array(neg_lefts) + np.array(neg_widths))
dataw = np.max(points) - np.min(points)

# draw invisible bars just for sizing the axes
label_padding = np.array([0.1*dataw if w < 1 else 0 for w in pos_widths])
plt.barh(pos_inds, np.array(pos_widths) + label_padding + 0.02*dataw,
        left=np.array(pos_lefts) - 0.01*dataw, color=colors.red_rgb, alpha=0
label_padding = np.array([-0.1*dataw if -w < 1 else 0 for w in neg_widths])
plt.barh(neg_inds, np.array(neg_widths) + label_padding - 0.02*dataw,
        left=np.array(neg_lefts) + 0.01*dataw, color=colors.blue_rgb, alpha=

# define variable we need for plotting the arrows
head_length = 0.08
bar_width = 0.8
xlen = plt.xlim()[1] - plt.xlim()[0]
fig = plt.gcf()
ax = plt.gca()
bbox = ax.get_window_extent().transformed(fig.dpi_scale_trans.inverted())
width = bbox.width
bbox_to_xscale = xlen/width
hl_scaled = bbox_to_xscale * head_length
renderer = fig.canvas.get_renderer()

# draw the positive arrows
for i in range(len(pos_inds)):
    dist = pos_widths[i]
    arrow_obj = plt.arrow(
        pos_lefts[i], pos_inds[i], max(dist-hl_scaled, 0.000001), 0,
        head_length=min(dist, hl_scaled),
        color=colors.red_rgb, width=bar_width,
        head_width=bar_width,
    )

```

```

if pos_low is not None and i < len(pos_low):
    plt.errorbar(
        pos_lefts[i] + pos_widths[i], pos_inds[i],
        xerr=np.array([[pos_widths[i] - pos_low[i]], [pos_high[i] - pos_w
        ecolor=colors.light_red_rgb,
    )

txt_obj = plt.text(
    pos_lefts[i] + 0.5*dist, pos_inds[i], format_value(pos_widths[i], '%+
    horizontalalignment='center', verticalalignment='center', color="white
    fontsize=12,
)
text_bbox = txt_obj.get_window_extent(renderer=renderer)
arrow_bbox = arrow_obj.get_window_extent(renderer=renderer)

# if the text overflows the arrow then draw it after the arrow
if text_bbox.width > arrow_bbox.width:
    txt_obj.remove()

    txt_obj = plt.text(
        pos_lefts[i] + (5/72)*bbox_to_xscale + dist, pos_inds[i], format_
        horizontalalignment='left', verticalalignment='center', color=col
        fontsize=12,
    )

# draw the negative arrows
for i in range(len(neg_inds)):
    dist = neg_widths[i]

    arrow_obj = plt.arrow(
        neg_lefts[i], neg_inds[i], -max(-dist-hl_scaled, 0.000001), 0,
        head_length=min(-dist, hl_scaled),
        color=colors.blue_rgb, width=bar_width,
        head_width=bar_width,
    )

    if neg_low is not None and i < len(neg_low):
        plt.errorbar(
            neg_lefts[i] + neg_widths[i], neg_inds[i],
            xerr=np.array([[neg_widths[i] - neg_low[i]], [neg_high[i] - neg_w
            ecolor=colors.light_blue_rgb,
        )

    txt_obj = plt.text(

```

```

        neg_lefts[i] + 0.5*dist, neg_inds[i], format_value(neg_widths[i], '%+1.2f'),
        horizontalalignment='center', verticalalignment='center', color="white",
        fontsize=12,
    )
    text_bbox = txt_obj.get_window_extent(renderer=renderer)
    arrow_bbox = arrow_obj.get_window_extent(renderer=renderer)

    # if the text overflows the arrow then draw it after the arrow
    if text_bbox.width > arrow_bbox.width:
        txt_obj.remove()

    txt_obj = plt.text(
        neg_lefts[i] - (5/72)*bbox_to_xscale + dist, neg_inds[i], format_value(neg_widths[i], '%+1.2f'),
        horizontalalignment='right', verticalalignment='center', color="white",
        fontsize=12,
    )

# draw the y-ticks twice, once in gray and then again with just the feature names
# The 1e-8 is so matplotlib 3.3 doesn't try and collapse the ticks
ytick_pos = list(range(num_features)) + list(np.arange(num_features)+1e-8)
plt.yticks(ytick_pos, yticklabels[:-1] + [label.split('=')[1] for label in yticklabels[-1:]])

# put horizontal lines for each feature row
for i in range(num_features):
    plt.axhline(i, color="#cccccc", lw=0.5, dashes=(1, 5), zorder=-1)

# mark the prior expected value and the model prediction
plt.axvline(base_values, 0, 1/num_features, color="#bbbbbb", linestyle="--", zorder=1)
fx = base_values + values.sum()
plt.axvline(fx, 0, 1, color="#bbbbbb", linestyle="--", linewidth=0.5, zorder=2)

# clean up the main axis
plt.gca().xaxis.set_ticks_position('bottom')
plt.gca().yaxis.set_ticks_position('none')
plt.gca().spines['right'].set_visible(False)
plt.gca().spines['top'].set_visible(False)
plt.gca().spines['left'].set_visible(False)
ax.tick_params(labelsize=13)
#plt.xlabel("\nModel output", fontsize=12)

#ax.set_xlim(0.00,1.00)

# draw the E[f(X)] tick mark
xmin, xmax = ax.get_xlim()
ax2 = ax.twinx()

```

```

ax2.set_xlim(xmin, xmax)
#ax2.set_xlim(0.00,1.00)
ax2.set_xticks([base_values, base_values+1e-8]) # The 1e-8 is so matplotlib
ax2.set_xticklabels(["\n$base$ $value$", "\n\t $ = " + format_value(base_value
ax2.spines['right'].set_visible(False)
ax2.spines['top'].set_visible(False)
ax2.spines['left'].set_visible(False)

# draw the f(x) tick mark
ax3 = ax2.twinx()
ax3.set_xlim(xmin, xmax)
#ax3.set_xlim(0.00,1.00)
# The 1e-8 is so matplotlib 3.3 doesn't try and collapse the ticks
ax3.set_xticks([base_values + values.sum(), base_values + values.sum() + 1e-8
ax3.set_xticklabels(["$prediction$", "\t\t $ = $" + "$" + format_value(fx, "%
tick_labels = ax3.xaxis.get_majorticklabels()
tick_labels[0].set_transform(tick_labels[0].get_transform(
) + matplotlib.transforms.ScaledTranslation(-10/72., 0, fig.dpi_scale_trans))
tick_labels[1].set_transform(tick_labels[1].get_transform(
) + matplotlib.transforms.ScaledTranslation(12/72., 0, fig.dpi_scale_trans))
tick_labels[1].set_color("#999999")
ax3.spines['right'].set_visible(False)
ax3.spines['top'].set_visible(False)
ax3.spines['left'].set_visible(False)

# adjust the position of the E[f(X)] = x.xx label
tick_labels = ax2.xaxis.get_majorticklabels()
tick_labels[0].set_transform(tick_labels[0].get_transform(
) + matplotlib.transforms.ScaledTranslation(-20/72., 0, fig.dpi_scale_trans))
tick_labels[1].set_transform(tick_labels[1].get_transform(
) + matplotlib.transforms.ScaledTranslation(22/72., -1/72., fig.dpi_scale_tra

tick_labels[1].set_color("#999999")

# color the y tick labels that have the feature values as gray
# (these fall behind the black ones with just the feature name)
tick_labels = ax.yaxis.get_majorticklabels()
for i in range(num_features):
    tick_labels[i].set_color("#999999")

if show:
    plt.show()
else:
    return plt.gca()
shap.plots.waterfall = patch_waterfall

```


Set up tutorial examples

Start by training the "should you bring an umbrella?" model

```
preX = pd.read_csv("Umbrella.csv")
preX = preX.sample(frac=1)
X_display = preX.iloc[:, :-1]
y_display = preX.iloc[:, -1]

PRECIPITATION = {
    "none": 0,
    "drizzle": 1,
    "rain": 2,
    "snow": 3,
    "sleet": 4,
    "hail": 5
}

y = y_display
X = X_display
X = X.replace({"Precipitation": PRECIPITATION})


X_train = X.iloc[:300]
y_train = y.iloc[:300]

X_test = X.iloc[300:]
y_test = y.iloc[300:]



d_train = lightgbm.Dataset(X_train, label=y_train)
d_test = lightgbm.Dataset(X_test, label=y_test)

params = {
    "max_bin": 512,
    "learning_rate": 0.05,
    "boosting_type": "gbdt",
    "objective": "binary",
    "metric": "binary_logloss",
    "num_leaves": 10,
    "verbose": -1,
    "min_data": 100,
    "boost_from_average": True,
    "keep_training_booster": True
}
```

```
#model = lgb.train(params, d_train, 10000, valid_sets=[d_test]) #early_stopping_rounds=100
model = lightgbm.LGBMClassifier(max_bin= 512,
                                learning_rate= 0.05,
                                boosting_type= "gbdt",
                                objective= "binary",
                                metric= "binary_logloss",
                                num_leaves= 10,
                                verbose= -1,
                                min_data= 100,
                                boost_from_average= True)
model.fit(X_train, y_train)
```

 <ipython-input-8-59731d8556ad>:17: FutureWarning: Downcasting behavior in `replace` is deprecated. In a future version, this will default to `False` to preserve dtype. Specify `downcast='integer'` to retain the current behavior. For more information, see [this issue](#).


```
X = X.replace({"Precipitation":PRECIPITATION})
```

 **LGBMClassifier** 

```
LGBMClassifier(boost_from_average=True, learning_rate=0.05, max_bin=512,
               metric='binary_logloss', min_data=100, num_leaves=10,
               objective='binary', verbose=-1)
```

Find the location of one of the two tutorial examples


```
print(X.loc[(X['Precipitation'] == 5) & (X['Temperature'] == 23) & (X['Wind(mph)' == 10)])
print(X.loc[(X['Precipitation'] == 0) & (X['Temperature'] == 70) & (X['Wind(mph)' == 30)])
theloc = X.index.get_loc(330)
```



	Precipitation	Temperature	Wind(mph)
330	5	23	10
	Precipitation	Temperature	Wind(mph)
96	0	70	30

Generate a tutorial explanation

```
explainer = shap.Explainer(model, X, model_output="probability")
#shap_values = explainer(X)
mynewexp = shap._explanation.Explanation(values=np.array([0.08, 0.26, -0.12]), ba
shap.plots.waterfall(mynewexp, max_display=20)
```



✓ Loan Instances

Edit and prepare dataset

```
# load dataset
```

```

X,y = shap.datasets.adult()
X_display,y_display = shap.datasets.adult(display=True)

EDUCATION_NUM = {
    16.0: "Doctorate",
    15.0: "Prof. School",
    14.0: "Masters",
    13.0: "Bachelors",
    12.0: "Some College",
    11.0: "Associate", #Assoc-acdm
    10.0: "Vocational", #Assoc-voc
    9.0: "HS grad",
    8.0: "12th",
    7.0: "11th",
    6.0: "10th",
    5.0: "9th",
    4.0: "7th-8th",
    3.0: "5th-6th",
    2.0: "1st-4th",
    1.0: "Preschool"
}

OCCUPATION_NUM = {
    "Tech-support": "Tech Support",
    "Craft-repair": "Craft/Repair",
    "Other-service": "Other Service",
    "Sales": "Sales",
    "Exec-managerial": "Exec. Managerial",
    "Prof-specialty": "Prof. Specialty",
    "Handlers-cleaners": "Handler/Cleaner",
    "Machine-op-inspct": "Machine Op. Inspector",
    "Adm-clerical": "Admin. Clerical",
    "Farming-fishing": "Farming/Fishing",
    "Transport-moving": "Transport/Moving",
    "Priv-house-serv": "Private House Service",
    "Protective-serv": "Protective Service",
    "Armed-Forces": "Armed Forces"
}

X_display = X_display.replace({"Education-Num":EDUCATION_NUM})
X_display = X_display.replace({"Occupation":OCCUPATION_NUM})
X = X.rename(columns={"Education-Num": "Education"})
X_display = X_display.rename(columns={"Education-Num": "Education"})#, "Hours per

X = X.drop(['Capital Loss', 'Capital Gain', 'Race', 'Relationship', 'Country', 'W

```

```
X_display = X_display.drop(['Capital Loss', 'Capital Gain', 'Race', 'Relationship'])

# create a train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
d_train = lgb.Dataset(X_train, label=y_train)
d_test = lgb.Dataset(X_test, label=y_test)
```

Train the model

```
params = {
    "max_bin": 512,
    "learning_rate": 0.05,
    "boosting_type": "gbdt",
    "objective": "binary",
    "metric": "binary_logloss",
    "num_leaves": 10,
    "verbose": -1,
    "min_data": 100,
    'objective': 'multi:softprob',
    "boost_from_average": True
}

params_xgb={
    'base_score':0.5,
    'learning_rate':0.05,
    'max_depth':5,
    'min_child_weight':100,
    'n_estimators':200,
    'num_class': 2,
    'nthread':-1,
    'objective': 'multi:softprob',
    'seed':2018,
    'eval_metric': 'auc'
}

model = lgb.LGBMClassifier(max_bin= 512,
    learning_rate= 0.05,
    boosting_type= "gbdt",
    objective= "binary",
    metric= "binary_logloss",
    num_leaves= 10,
    verbose= -1,
    min_data= 100,
    boost_from_average= True)
model.fit(X_train, y_train)
```

**LGBMClassifier**

```
LGBMClassifier(boost_from_average=True, learning_rate=0.05, max_bin=512,
    metric='binary_logloss', min_data=100, num_leaves=10,
    objective='binary', verbose=-1)
```

Our 7 loan application instances

```
#val = 610 # Woman Side-by-side
#val = 11116 # Man Side-by-side
#val = 32353 # Man 3
#val = 217 # Man 2
#val = 15040 # Man 1
#val = 32429 # Woman 3
val = 32556 # Woman 2
#val = 91#91 # Woman 1

theloc = val
```

Generate SHAP Explanation

```
explainer = shap.Explainer(model, X, model_output="probability")
shap_values = explainer(X)
```

🔄 100%|=====| 32418/32561 [01:47<00:00]

```
#shap_values_standin0 = pd.Series({'Age': 0.0307, 'Education': -0.0287, 'Occupation': 0.0307})
shap_values_standin0 = pd.Series({'Age': -0.14, 'Education': 0.0416, 'Occupation': 0.0307})
#shap_values_standin0 = pd.Series({'Age': 0.1209, 'Education': 0.3008, 'Occupation': 0.0307})
#shap_values_standin0 = pd.Series({'Age': -0.2119, 'Education': 0.0011, 'Occupation': 0.0307})
#shap_values_standin0 = pd.Series({'Age': 0.0565, 'Education': 0.1427, 'Occupation': 0.0307})
#shap_values_standin0 = pd.Series({'Age': -0.0012, 'Education': -0.189, 'Occupation': 0.0307})
#shap_values_standin0 = pd.Series({'Age': 0.0774, 'Education': 0.1962, 'Occupation': 0.0307})
#shap_values_standin0 = pd.Series({'Age': 0.0668, 'Education': 0.1619, 'Occupation': 0.0307})

mynewexp = shap._explanation.Explanation(values=shap_values_standin0, base_values = 0.0307)

shap.plots.waterfall(mynewexp, max_display=20)
```

```
<ipython-input-7-465cce54bdf2>:118: FutureWarning: Series.__getitem__ treating  
      sval = values[order[i]]  
<ipython-input-7-465cce54bdf2>:140: FutureWarning: Series.__getitem__ treating  
      if np.issubdtype(type(features[order[i]]), np.number):  
<ipython-input-7-465cce54bdf2>:141: FutureWarning: Series.__getitem__ treating  
      yticklabels[rng[i]] = format_value(float(features[order[i]]), "%0.03f") + "  
<ipython-input-7-465cce54bdf2>:143: FutureWarning: Series.__getitem__ treating  
      yticklabels[rng[i]] = str(features[order[i]]) + " = " + str(feature_names[or  
      prediction = 0.078
```

